

SOAR Service Mesh

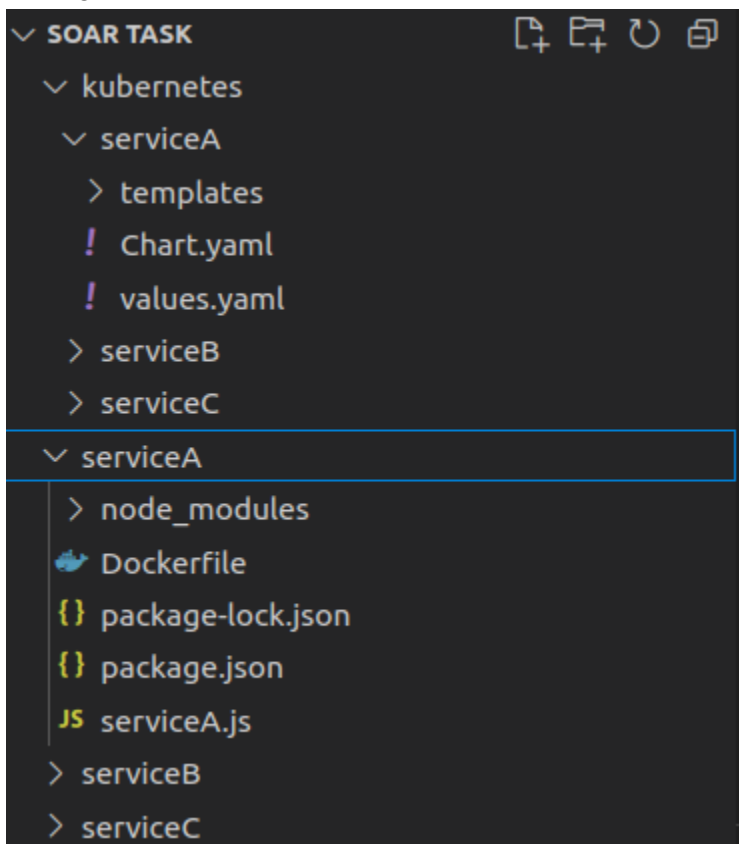
Deploying Micorservices	2
Installing linkerd	5
Prometheus and Grafanna	8

Deploying Micorservices

Creating a namespace soar-task

```
(base) ahmed@ahmed-Inspiron-5491-2n1:~/Desktop/soar task$ kubectl get namespaces
NAME                STATUS    AGE
arms-prom           Active    3d2h
default             Active    3d2h
dev                 Active    3d
kube-node-lease     Active    3d2h
kube-public         Active    3d2h
kube-system         Active    3d2h
soar-task           Active    8s
```

Setting the folder structure we will assume a monorepo for simplicity's sake



Kubernetes : helm packages for the microserivces

Adding express docker files

```
FROM node:18-alpine

WORKDIR /usr/src/app

COPY package*.json ./

RUN npm ci --verbose

COPY . .

CMD ["node", "serviceA.js"]
```

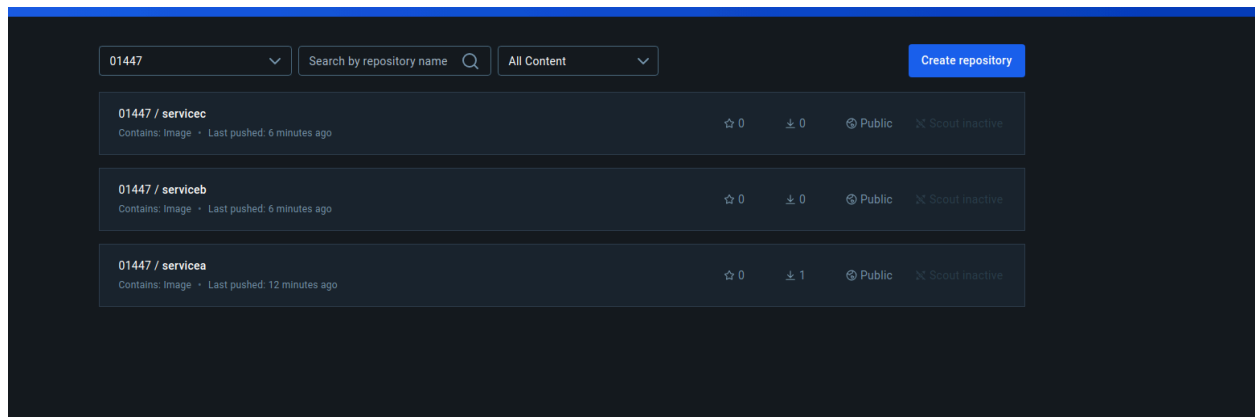
Node18-alpine light weight version of node

Setting the working directory to /usr/src/app

Copying packages first then copying the whole dir as the source code is latest thing that would change

Npm ci is clean installation making sure package.json is maintainable in prod

Build all images and push them to dockerhub



After Adding helm packages

Each service has been added to its own helm package for modularity.

Ex:

```
helm upgrade servicec ./ -n soar-task --install
```

```
(base) ahmed@ahmed-Inspiron-5491-2n1:~/Desktop/soar task/kubernetes/serviceC$ kubectl get pods -n soar-task
```

NAME	READY	STATUS	RESTARTS	AGE
service-a-servicea-656d9d7b78-48scg	1/1	Running	0	9m51s
service-b-serviceb-85fbb997df-7pscj	1/1	Running	0	25s
service-c-servicec-cff5d99fb-qp6bk	1/1	Running	0	8s

Since we still did not install linkerd it will only be 1/1 replica pods

Installing linkerd

Following the [docs](#) there are multiple sources to install from their opensource docs is the most informative :

Linkerd deployment models:

Host based and sidecar proxy based since wer deploying dockerized microservices on kubernetes we will be using a sidecar proxy

```
NAME                READY    UP-TO-DATE    AVAILABLE    AGE
linkerd-destination  0/1      1              0            5m
linkerd-identity     1/1      1              1            5m3s
linkerd-proxy-injector 1/1      1              1            4m59s
(base) ahmed@ahmed-Inspiron-5491-2n1: ~/Desktop/soar task$
```

Injecting the sidecar:

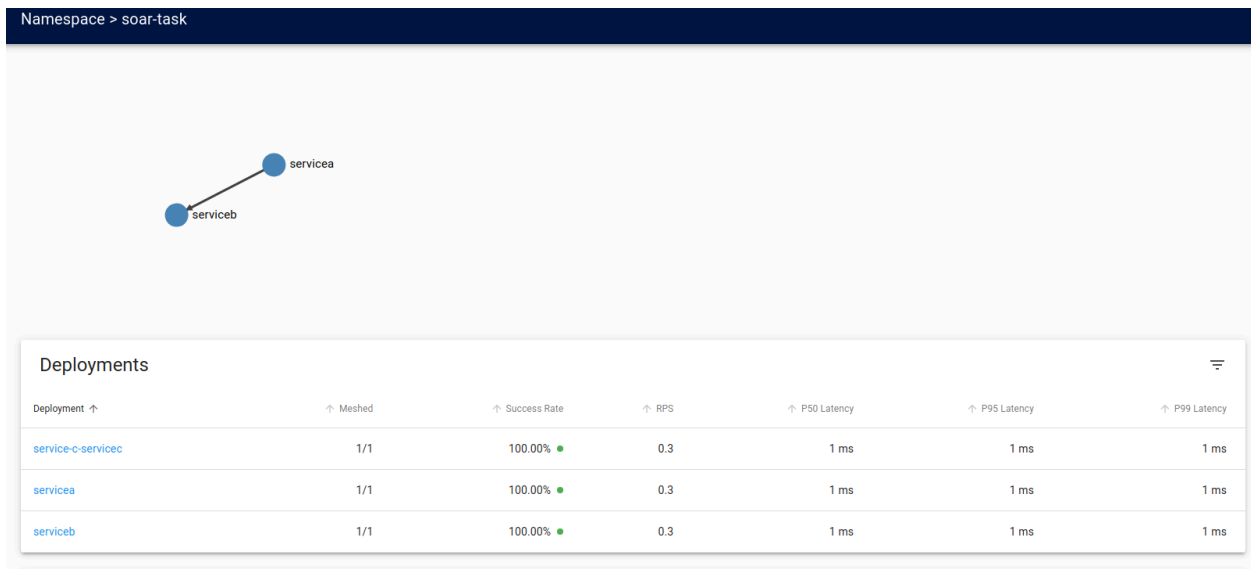
Either imperative or declarative inside the kuebrentes manifest we went with the adding it inside the kubernetes manifest in the following snippet

```
template:
  metadata:
    labels:
      app: {{ include "servicec.name" . }}
      release: {{ .Release.Name }}
    annotations:
      linkerd.io/inject: enabled
```

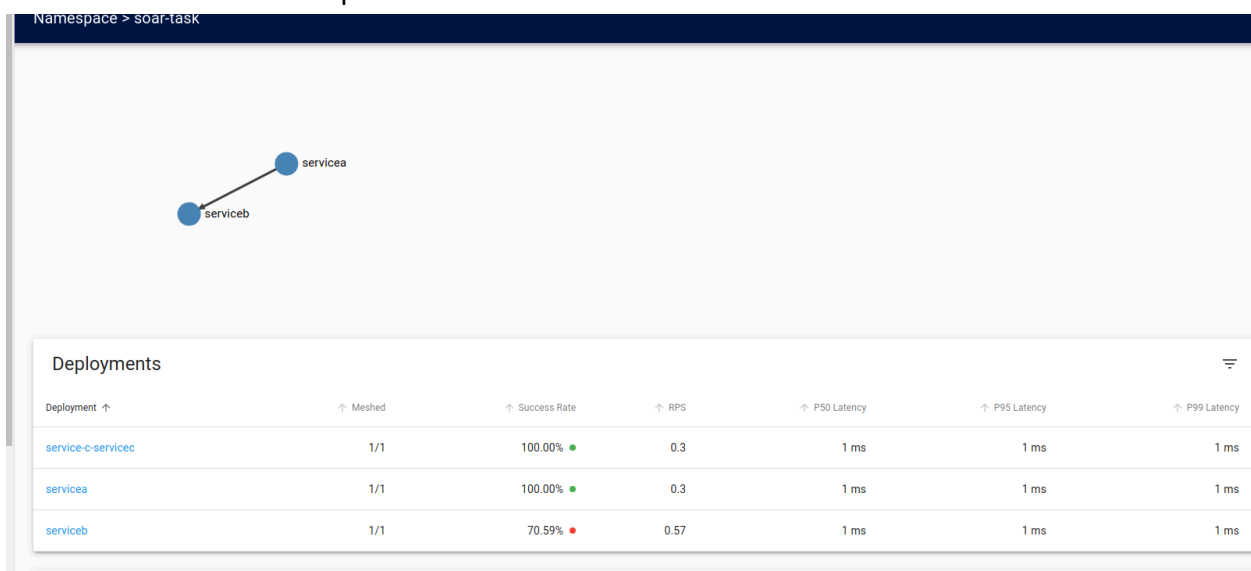
After injecting the sidecar:

```
• (base) ahmed@ahmed-Inspiron-5491-2n1:~/Desktop/soar task/kubernetes/serviceC$ kubectl get pods -n soar-task
NAME                READY    STATUS      RESTARTS    AGE
servicea-57ccbfddcd-wmmff  2/2      Running      0            68m
serviceb-6d746df9b9-8r5bs  2/2      Running      0            18m
servicec-6647bd89bd-7wmvr  2/2      Running      0            32s
serviced-78d744879f-qrbqg  1/1      Terminating 0            4m47s
```

At a first glance all services seem to be working which doesn't seem to be making sense the reason is serviceB's error is handled in a try catch statement



We'll change microserviceB so instead of handling the error it can sometimes throw a 500 error and we will send a curl request from servicea to serviceb:



It reflects right away !

So for serviceB using error handling would do the trick

Now that we finally got through with service B lets check serviceC the high latency one when we start hitting its endpoint we can see it reflects in the latency p50, p95 and p99 latencies all represent the percentiles of different latencies so under normal circumstances if the total number of request latencies is normally distributed our p99 latency would mean that around 99 percent of the requests hitting our microservice would be 99 percent and same applies to p95 and p50.

FROM

[deploy/servicea](#)

GET

/

2

10.0 s

10.0 s

10.0 s

100.00%

Current Top query

linkerd viz top deployment/servicec --namespace soar-task

Inbound

Resource ↑	↑ Meshed	↑ Success Rate	↑ RPS	↑ P50 Latency	↑ P95 Latency	↑ P99 Latency
deploy/prometheus	1/1	100.00% 	0.1	2 ms	2 ms	2 ms
deploy/servicea	1/1	100.00% 	0.03	15.0 s	19.5 s	19.9 s
deploy/tap	1/1	---	---	---	---	---

Pods



Reflection: since its a generic timeout function it is apparent that latency in response times would drop once we reduce the time in the function however in production environments it could be multiple factors such as the cpu resources of our kubernetes nodes, the allowed cpu the microservice can take up from the node and finally the source code itself where it might be slow from the data layer maybe slow query performance or the endpoint calling multiple functions that could be better broken down.

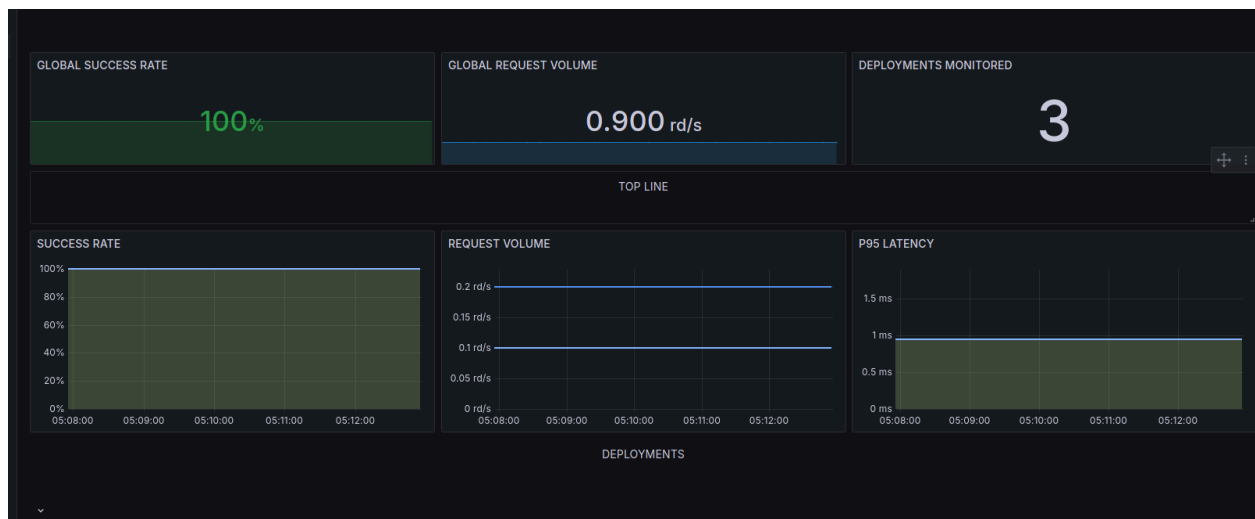
Prometheus and Grafanna

Now that we've finished all of our service mesh setup and put down our reflection we need one final integration along with linkerd inorder to persist the metrics getting stored.

Linkerd viz extension that we've installed installs prometheus so our final integration is to install [grafana](#).

After the follow along get a grafana icon next to linkerd:

linkerd	3/3	100.00% 	2.52	1 ms	1 ms	2 ms	
linkerd-viz	5/5	100.00% 	3.3	1 ms	2 ms	4 ms	
soar-task	3/3	100.00% 	0.9	1 ms	1 ms	1 ms	



Once we click on a deployment we get the full grafana dashboard with the linkerd metrics already present.